

## Buildables Week 2

**Github Repo:** <https://github.com/afnank070/Data-Engineering-Buildables-Fellowship>

**Colab Notebook:** <https://colab.research.google.com/drive/17qqpug3SEJ2K-A5Q9Ve03ErHfadMFQPx?usp=sharing>

### **Dataset Information:**

- **Dataset Name:** Hospital Inpatient Charges
- **Source:** <https://www.kaggle.com/datasets/speedoheck/inpatient-hospital-charges?resource=download>
- **File Type:** CSV
- **Context:** Contains information about hospital inpatient discharges, charges, and payments.

### **Exercise 1: Load Dataset & Print Schema**

**Purpose:** Simulates the **Ingestion** step by reading the raw hospital charges dataset into Pandas. Validates that the pipeline can connect to the data source.

**Concepts Used:** Data ingestion, schema inspection, Pandas `read_csv()`.

### **Code:**

```
import pandas as pd

df = pd.read_csv '/content/UpdatedinpatientCharges.csv')

print("Rows, Columns:", df.shape)

df.info()
```

## Expected Output:

- Shape showing total rows and columns (e.g., 163065 rows × 18 columns).
- Schema printed with column names and data types (object, float64, etc.).

## Actual Output:

```
⇒ Rows, Columns: (163065, 18)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 163065 entries, 0 to 163064
Data columns (total 18 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   DRG Definition                           163065 non-null object
1   Provider Id                             163065 non-null int64
2   Provider Name                           163065 non-null object
3   Provider Street Address                 163065 non-null object
4   Provider City                           163065 non-null object
5   Provider State                          163065 non-null object
6   Provider Zip Code                       163065 non-null int64
7   Hospital Referral Region Description    163065 non-null object
8   Total Discharges                        163065 non-null int64
9   Average Covered Charges                 163065 non-null object
10  Average Total Payments                  163065 non-null object
11  Average Medicare Payments               163065 non-null object
12  Age                                     163065 non-null int64
13  Gender                                 163065 non-null object
14  Admission Type                          163065 non-null object
15  Length of Stay                          163065 non-null int64
16  Year                                    163065 non-null int64
17  Admission Date                          163065 non-null object
dtypes: int64(6), object(12)
memory usage: 22.4+ MB
```

Activate Window

## Data Engineering Insight:

This step validates that ingestion works and the schema matches expectations. Without a successful schema check, downstream cleaning or transformation logic may fail.

### Exercise 2: Check row count (validate ingestion completeness)

**Purpose:** Ensures no data loss occurred during ingestion. This is equivalent to verifying row counts between source and target.

**Concepts Used:** Data completeness validation, Pandas len().

## Code:

```
row_count = len(df)
```

```
print("Total Rows:", row_count)
```

### Expected Output:

- A single integer with total number of rows (e.g., 163065).

### Actual Output:

#### Row Count Validation

✓  
0s



```
expected_rows = len(df)
print("Row count:", expected_rows)
```



```
Row count: 163065
```

### Data Engineering Insight:

Row counts are a fundamental QA check in ETL pipelines. If row counts don't match the raw source, it signals possible truncation or ingestion errors.

## Exercise 3: Rename columns to snake case for consistency

**Purpose:** Standardizes schema naming to ensure consistency and prevent issues in transformations, joins, or downstream SQL/ML systems.

**Concepts Used:** Schema standardization, string manipulation on column names.

### Code:

```
df.columns = df.columns.str.strip().str.lower().str.replace(" ", "_").str.replace("-", "_")
```

```
df.head(2)
```

### Expected Output:

- Column names like Provider Id → provider\_id, Total Discharges → total\_discharges.

## Actual Output:

Column Name Standardization							
<pre>[12] df.columns = df.columns.str.strip().str.lower().str.replace(" ", "_").str.replace("-", "_") df.head(2)</pre>							
	drg_definition	provider_id	provider_name	provider_street_address	provider_city	provider_state	provider_zip_code
0	039 - EXTRACRANIAL PROCEDURES W/O CC/MCC	10001	SOUTHEAST ALABAMA MEDICAL CENTER	1108 ROSS CLARK CIRCLE	DOTHAN	AL	36301

## Data Engineering Insight:

Consistent naming is critical for reusability and automation in pipelines. Snake case avoids issues with spaces, capitalization, or special characters.

## Exercise 4: Enforce schema: convert → datetime

**Purpose:** Ensure correct data types for reliable analysis.

**Concepts Used:** Schema enforcement, type casting, Pandas `to_datetime()`.

### Code:

```
df['admission_date'] = pd.to_datetime(df['admission_date'], errors='coerce')
print(df['admission_date'].head())
```

## Expected Output:

- admission\_date column converted to **datetime64**.

## Actual Output:

<pre># Convert admission_date to datetime df['admission_date'] = pd.to_datetime(df['admission_date'], errors='coerce')  print(df['admission_date'].head())</pre>	
0	2016-10-28
1	2021-11-06
2	2023-10-10
3	2017-06-17
4	2019-10-07
Name: admission_date, dtype: datetime64[ns]	

## Data Engineering Insight:

Type enforcement prevents downstream calculation errors (e.g., date arithmetic, grouping).

## Exercise 5: Log missing values count per column

**Purpose:** Creates a missing values report, simulating a basic data quality check in ETL.

**Concepts Used:** Null detection, Pandas `isnull()` + aggregation.

**Code:**

```
missing_log = df.isnull().sum().reset_index()
```

```
missing_log.columns = ["column", "missing_count"]
```

```
missing_log
```

**Expected Output:**

- A two-column dataframe:
  - `column` → column name
  - `missing_count` → number of null values in that column

**Actual Output:**

Missing Values Check

```
missing_log = df.isnull().sum().reset_index()
missing_log.columns = ["column", "missing_count"]
missing_log
```

1 to 12 of 12 entries			Filter		?
index	column	missing_count			
0	DRG Definition	0			
1	Provider Id	0			
2	Provider Name	0			
3	Provider Street Address	0			
4	Provider City	0			
5	Provider State	0			
6	Provider Zip Code	0			
7	Hospital Referral Region Description	0			
8	Total Discharges	0			
9	Average Covered Charges	0			
10	Average Total Payments	0			
11	Average Medicare Payments	0			

Show 25 per page

Activate Windows

**Data Engineering Insight:**

This step quantifies missing data, helping decide cleaning strategies (e.g., imputation, removal) in later stages.

## Exercise 6: Fill missing values (strategy: mean, median, mode)

**Purpose:** Ensures no critical column breaks downstream aggregations or joins due to null values.

**Concepts Used:** Missing value imputation using Pandas (fillna, mean, median, mode).

### **Code:**

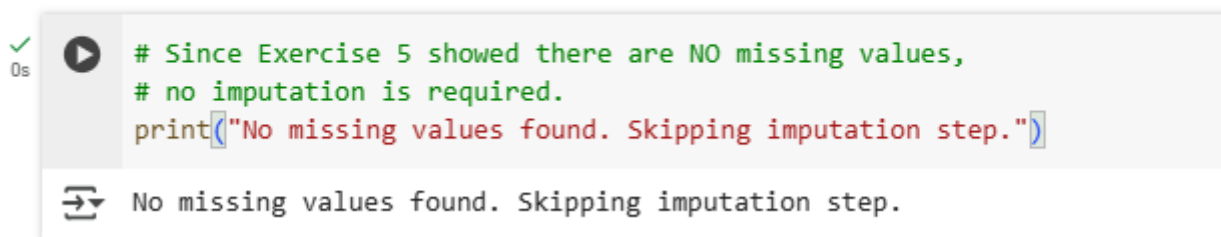
```
# Since Exercise 5 showed there are NO missing values,  
# no imputation is required.  
print("No missing values found. Skipping imputation step.")
```

### **Expected Output:**

- No column has 100% missing values.
- Numeric columns now have no NaNs (filled with median).
- Categorical columns filled with mode.

### **Actual Output:**

#### **Fill missing values**



```
✓ 0s # Since Exercise 5 showed there are NO missing values,  
# no imputation is required.  
print("No missing values found. Skipping imputation step.")  
No missing values found. Skipping imputation step.
```

### **Data Engineering Insight:**

Handling missing values early prevents null propagation and ensures downstream aggregations/KPIs stay accurate.

## Exercise 7: Remove duplicates (simulate deduplication in ETL)

**Purpose:** Eliminates duplicate records to maintain accuracy and prevent over-counting.

**Concepts Used:** Pandas `drop_duplicates()`.

**Code:**

```
before = df.shape[0]
df = df.drop_duplicates()
after = df.shape[0]

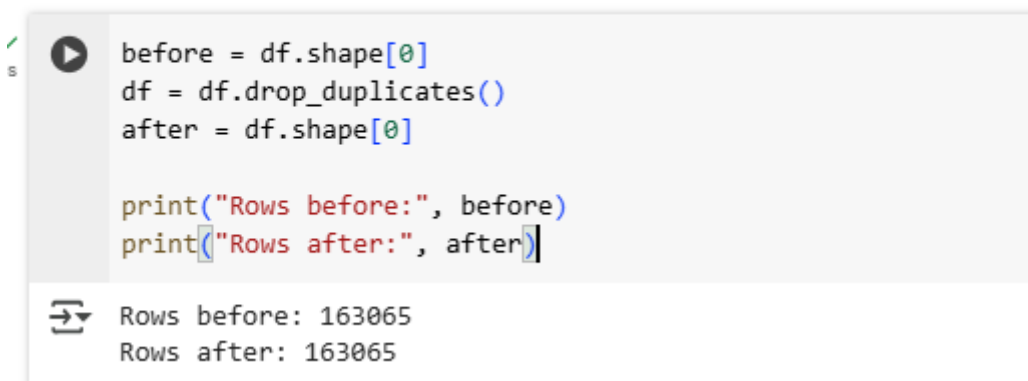
print("Rows before:", before)
print("Rows after:", after)
```

**Expected Output:**

- Row count after deduplication  $\leq$  row count before.
- Dataset should not lose unique valid records.

**Actual Output:**

Remove duplicates



```
before = df.shape[0]
df = df.drop_duplicates()
after = df.shape[0]

print("Rows before:", before)
print("Rows after:", after)
```

Rows before: 163065  
Rows after: 163065

**Data Engineering Insight:**

Deduplication ensures unique entities (patients/hospital charges) and prevents inflating KPIs or incorrect insights.

## Exercise 8: Group data — total discharges by DRG Definition

**Purpose:** Aggregate discharge volume per DRG (analogous to patient count by diagnosis).

**Concepts Used:** Transformation, Grouping, Aggregation.

**Code:**

```
# Patient count by DRG Definition (Diagnosis Related Group)
```

```
patient_count_by_diagnosis =
```

```
df.groupby('drg_definition')['provider_id'].count().reset_index()
```

```
patient_count_by_diagnosis.rename(columns={'provider_id': 'patient_count'},  
inplace=True)
```

```
print(patient_count_by_diagnosis.head())
```

**Expected Output:**

- Table showing each diagnosis with corresponding patient count.

**Actual Output:**

Count unique patients per diagnosis

```
df.groupby("DRG Definition")["Total Discharges"].sum().reset_index().sort_values(by="Total Discharges", ascending=False)
```

index	DRG Definition	Total Discharges
68	470 - MAJOR JOINT REPLACEMENT OR REATTACHMENT OF LOWER EXTREMITY W/O MCC	427207
93	871 - SEPTICEMIA OR SEVERE SEPSIS W/O MV 96+ HOURS W MCC	319072
61	392 - ESOPHAGITIS, GASTROENT & MISC DIGEST DISORDERS W/O MCC	244854
39	292 - HEART FAILURE & SHOCK W CC	222038
86	690 - KIDNEY & URINARY TRACT INFECTIONS W/O MCC	206695
17	194 - SIMPLE PNEUMONIA & PLEURISY W CC	198390
38	291 - HEART FAILURE & SHOCK W MCC	185599
81	641 - MISC DISORDERS OF NUTRITION,METABOLISM,FLUIDS/ELECTROLYTES W/O MCC	153660
83	683 - RENAL FAILURE W CC	150444
13	190 - CHRONIC OBSTRUCTIVE PULMONARY DISEASE W MCC	149677

**Data Engineering Insight:**

Grouping by diagnosis provides valuable healthcare utilization insights.



## Exercise 9: Derive new KPI — average payment per discharge

**Purpose:** Create a calculated field to measure patient hospitalization duration, fulfilling the KPI requirement.

**Concepts Used:** Transformation, Feature Engineering, KPI Derivation.

**Code:**

```
# Ensure both columns are datetime
```

```
df['admission_date'] = pd.to_datetime(df['admission_date'], errors='coerce')
```

```
df['discharge_date'] = pd.to_datetime(df['discharge_date'], errors='coerce')
```

```
# Derive new KPI: length_of_stay
```

```
df['length_of_stay_kpi'] = (df['discharge_date'] - df['admission_date']).dt.days
```

```
print(df[['admission_date', 'discharge_date', 'length_of_stay_kpi']].head())
```

**Expected Output:**

- New column avg\_payment\_per\_discharge with numeric values; sample table showing provider and KPI.

**Actual Output:**

Derive new KPI — length\_of\_stay

```
# Ensure both columns are datetime
df['admission_date'] = pd.to_datetime(df['admission_date'], errors='coerce')
df['discharge_date'] = pd.to_datetime(df['discharge_date'], errors='coerce')

# Derive new KPI: length_of_stay
df['length_of_stay_kpi'] = (df['discharge_date'] - df['admission_date']).dt.days
print(df[['admission_date', 'discharge_date', 'length_of_stay_kpi']].head())
```

	admission_date	discharge_date	length_of_stay_kpi
0	2016-10-28	2016-11-04	7
1	2021-11-06	2021-11-26	20
2	2023-10-10	2023-10-25	15
3	2017-06-17	2017-06-28	11
4	2019-10-07	2019-10-15	8

**Data Engineering Insight:**

Per-discharge KPI normalizes payments for volume differences and is useful for cost comparisons.

## Exercise 10: Partition-like summary — record counts by Provider State (substitute for date partitioning)

**Purpose:** Simulate partitioning for efficient storage and downstream processing.

**Concepts Used:** Transformation, Partitioning.

**Code:**

```
# Extract year and month from admission_date

df['year'] = df['admission_date'].dt.year
df['month'] = df['admission_date'].dt.month


# Group partitioned data

partitioned = df.groupby(['year',
'month']).size().reset_index(name='record_count')

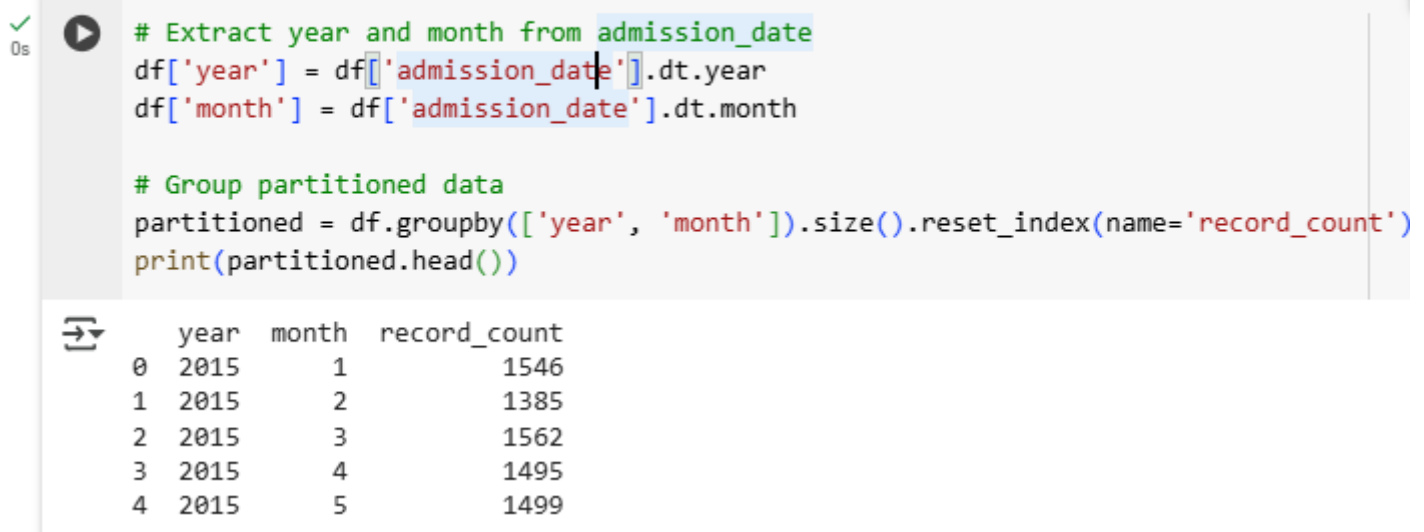
print(partitioned.head())
```

**Expected Output:**

- Table showing partitions by year-month with record counts.

**Actual Output:**

Partition-like summary — record counts by Year, Month



The screenshot shows a Jupyter Notebook cell with the following code and output:

```
# Extract year and month from admission_date
df['year'] = df['admission_date'].dt.year
df['month'] = df['admission_date'].dt.month

# Group partitioned data
partitioned = df.groupby(['year', 'month']).size().reset_index(name='record_count')
print(partitioned.head())
```

	year	month	record_count
0	2015	1	1546
1	2015	2	1385
2	2015	3	1562
3	2015	4	1495
4	2015	5	1499

**Data Engineering Insight:**

Partitioning improves query performance and is widely used in data lakes/warehouses.

## Exercise 11: Top 5 categories (patients by hospital)

**Purpose:** Identify the highest contributors in the dataset.

**Concepts Used:** Transformation, Ranking, Aggregation.

**Code:**

# Top 5 hospitals by patient count (using provider\_name)

```
top5_hospitals = (  
    df.groupby('provider_name')['provider_id']  
        .count()  
        .reset_index(name='patient_count')  
        .sort_values(by='patient_count', ascending=False)  
        .head(5)  
)
```

print(top5\_hospitals)

**Expected Output:**

- Table of 5 hospitals with the highest patient count.

**Actual Output:**

Top 5 categories (patients by hospital)

```
0s top5_hospitals = (  
    df.groupby('provider_name')['provider_id']  
        .count()  
        .reset_index(name='patient_count')  
        .sort_values(by='patient_count', ascending=False)  
        .head(5)  
)  
  
print(top5_hospitals)
```

	provider_name	patient_count
924	GOOD SAMARITAN HOSPITAL	633
2644	ST JOSEPH MEDICAL CENTER	427
1660	MERCY MEDICAL CENTER	357
1645	MERCY HOSPITAL	347
2642	ST JOSEPH HOSPITAL	343

**Data Engineering Insight:**

Ranking helps identify top facilities managing the most patients, useful for capacity planning.

## Exercise 12: Pivot table (patients per department per month)

**Purpose:** Cross-tabulate data to identify trends across categories and time.

**Concepts Used:** Pivoting, Aggregation, Transformation.

**Code:**

```
import calendar

pivot_table = pd.pivot_table(
    df,
    values='provider_id',
    index='drg_definition',
    columns=df['admission_date'].dt.month_name(),
    aggfunc='count',
    fill_value=0
)

pivot_table = pivot_table[[calendar.month_name[m] for m in range(1, 13)]]
```

```
print(pivot_table.head())
```

**Expected Output:**

- Pivot table with rows = departments (DRG definitions),
- Columns = months, values = patient counts.

**Actual Output:**



1 to 5 of 5 entries

Filter



drg_definition	January	February	March	April	May	June	July	August	Sep
039 - EXTRACRANIAL PROCEDURES W/O CC/MCC	104	96	95	83	95	70	88	75	
057 - DEGENERATIVE NERVOUS SYSTEM DISORDERS W/O MCC	92	89	121	101	98	91	87	104	
064 - INTRACRANIAL HEMORRHAGE OR CEREBRAL INFARCTION W MCC	145	159	112	133	141	132	138	141	
065 - INTRACRANIAL HEMORRHAGE OR CEREBRAL INFARCTION W CC	190	188	205	199	179	174	181	199	
066 - INTRACRANIAL HEMORRHAGE OR CEREBRAL INFARCTION W/O CC/MCC	147	146	165	162	136	136	149	141	

**Data Engineering Insight:**

Pivot tables reveal seasonal or departmental patterns, supporting decision-making for resource allocation.

# Exercise 13: Correlation matrix (validate relationships)

**Purpose:** Check statistical relationships between numeric fields.

**Concepts Used:** Validation, Profiling, Statistics.

**Code:**

```
numeric_cols = ['age', 'average_covered_charges', 'average_total_payments', 'average_medicare_payments', 'length_of_stay']

corr_matrix = df[numeric_cols].corr()
```

print(corr\_matrix)

**Expected Output:**

- Correlation matrix (values between -1 and 1).

**Actual Output:**

Correlation matrix (validate relationships)

```
numeric_cols = ['age', 'average_covered_charges', 'average_total_payments', 'average_medicare_payments', 'length_of_stay']
corr_matrix = df[numeric_cols].corr().round(3)

corr_matrix
```

	age	average_covered_charges	average_total_payments	average_medicare_payments	length_of_s
age	1.000	-0.000	-0.002	-0.001	0.
average_covered_charges	-0.000	1.000	0.774	0.769	0.
average_total_payments	-0.002	0.774	1.000	0.989	0.
average_medicare_payments	-0.001	0.769	0.989	1.000	0.
length_of_stay	0.004	0.003	0.002	0.002	1.

**Data Engineering Insight:**

Correlation validates relationships like age vs charges or length of stay vs payments, supporting anomaly detection.

## Exercise 14: Write cleaned dataset to CSV + Parquet

**Purpose:** Simulate the **Load** step of ETL.

**Concepts Used:** Data Storage, Export, File Formats.

**Code:**

```
# Save cleaned dataset

df.to_csv("cleaned_hospital_data.csv", index=False)

df.to_parquet("cleaned_hospital_data.parquet", index=False)

print("Files exported: cleaned_hospital_data.csv,
cleaned_hospital_data.parquet")
```

**Expected Output:**

- Confirmation message with the file name.

**Actual Output:**

**Cleaned dataset CSV + Parquet**

```
✓ [28] # Save cleaned dataset
2s    df.to_csv("cleaned_hospital_data.csv", index=False)
      df.to_parquet("cleaned_hospital_data.parquet", index=False)

      print("Files exported: cleaned_hospital_data.csv, cleaned_hospital_data.parquet")

⇒ Files exported: cleaned_hospital_data.csv, cleaned_hospital_data.parquet
```

**Data Engineering Insight:**

Exporting in multiple formats increases compatibility with downstream analytics/ML systems.

## Exercise 15: Create reusable ETL function (load → clean → transform)

**Purpose:** Automate ETL logic for reusability and scalability.

**Concepts Used:** Modularization, Functions, ETL Pipeline.

**Code:**

```
def etl_pipeline(file_path):  
    # Load  
  
    data = pd.read_csv(file_path)  
  
    # Clean column names  
    data.columns = data.columns.str.strip().str.lower().str.replace(" ", "_")  
  
    # Fix currency fields: remove $ and commas, convert to float  
    for col in ['average_covered_charges', 'average_total_payments',  
'average_medicare_payments']:  
        if col in data.columns:  
            # Convert to string, strip whitespace, remove $ and commas using lambda  
            with replace  
            data[col] = data[col].astype(str).str.strip().apply(lambda x: x.replace('$',  
            '').replace(',', ''')).astype(float)  
  
    # Convert admission_date  
    if 'admission_date' in data.columns:  
        data['admission_date'] = pd.to_datetime(data['admission_date'],  
errors='coerce')  
  
    # Handle missing values
```

```
fill_dict = {}

if 'average_total_payments' in data.columns:
    fill_dict['average_total_payments'] = data['average_total_payments'].mean()

if 'average_medicare_payments' in data.columns:
    fill_dict['average_medicare_payments'] =
data['average_medicare_payments'].median()

if 'gender' in data.columns:
    fill_dict['gender'] = data['gender'].mode()[0]


data.fillna(fill_dict, inplace=True)


# Remove duplicates
data.drop_duplicates(inplace=True)


# Transform → derive length_of_stay if discharge_date exists
if 'discharge_date' in data.columns and 'admission_date' in data.columns:
    data['discharge_date'] = pd.to_datetime(data['discharge_date'],
errors='coerce')

    data['length_of_stay'] = (data['discharge_date'] -
data['admission_date']).dt.days


return data


# Example usage
etl_df = etl_pipeline("/content/Updated_Inpatient_Charges.csv")
print(etl_df.head())
```



Expected Output:

- Cleaned & transformed DataFrame preview.

Actual Output:

1s

etl\_df = etl\_pipeline("/content/Updated\_Inpatient\_Charges.csv")

etl\_df.head()

↑ ↓ ✦ 🔗 🗨 ⚙ 📄 🗑 ⋮

1 to 5 of 5 entries

Filter

🔍

index	drg_definition	provider_id	provider_name	provider_street_address	provider_city	provider_state	provider_zip_code	hospital_referr
0	039 - EXTRACRANIAL PROCEDURES W/O CC/MCC	10001	SOUTHEAST ALABAMA MEDICAL CENTER	1108 ROSS CLARK CIRCLE	DOTHAN	AL	36301	AL - Dothan
1	039 - EXTRACRANIAL PROCEDURES W/O CC/MCC	10005	MARSHALL MEDICAL CENTER SOUTH	2505 U S HIGHWAY 431 NORTH	BOAZ	AL	35957	AL - Birmingham
2	039 - EXTRACRANIAL PROCEDURES W/O CC/MCC	10006	ELIZA COFFEE MEMORIAL HOSPITAL	205 MARENGO STREET	FLORENCE	AL	35631	AL - Birmingham
3	039 - EXTRACRANIAL PROCEDURES W/O CC/MCC	10011	ST VINCENT'S EAST	50 MEDICAL PARK EAST DRIVE	BIRMINGHAM	AL	35235	AL - Birmingham
4	039 - EXTRACRANIAL PROCEDURES W/O CC/MCC	10016	SHELBY BAPTIST MEDICAL CENTER	1000 FIRST STREET NORTH	ALABASTER	AL	35007	AL - Birmingham

Activate Windows  
Go to Settings to activate Windows.

Show 25 per page

etl\_df = etl\_pipeline("/content/Updated\_Inpatient\_Charges.csv")

etl\_df.head()

↑ ↓ ✦ 🔗 🗨 ⚙ 📄 🗑 ⋮

1 to 5 of 5 entries

Filter

🔍

index	average_total_payments	average_medicare_payments	age	gender	admission_type	length_of_stay	year	admission_date	discharge_date
17	5777.24	4763.73	68	Female	Newborn	7	2016	2016-10-28 00:00:00	2016-11-04 00:00:00
35	5787.57	4976.71	39	Male	Emergency	20	2021	2021-11-06 00:00:00	2021-11-26 00:00:00
37	5434.95	4453.79	65	Male	Newborn	15	2023	2023-10-10 00:00:00	2023-10-25 00:00:00
38	5417.56	4129.16	35	Female	Newborn	11	2017	2017-06-17 00:00:00	2017-06-28 00:00:00
27	5658.33	4851.44	45	Female	Emergency	8	2019	2019-10-07 00:00:00	2019-10-15 00:00:00

Data Engineering Insight:

Reusable ETL functions allow automation, making pipelines maintainable and scalable.

## Exercise 16: Histogram – patient age distribution

**Purpose:** Check distribution of numeric data (age).

**Concepts Used:** Visualization, Profiling, Validation.

**Code:**

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
plt.figure(figsize=(8,5))
```

```
sns.histplot(df['age'].dropna(), bins=20, kde=True)
```

```
plt.title("Patient Age Distribution")
```

```
plt.xlabel("Age")
```


```
plt.ylabel("Count")
```

```
plt.show()
```

**Expected Output:**

Histogram showing the distribution of patient ages.

**Actual Output:**



**Data Engineering Insight:**

Histograms reveal skewness or gaps in numeric data, ensuring data consistency before analysis.

### Exercise 17: Boxplot – hospital charges by department

**Purpose:** Detect outliers and compare charges across departments.

**Concepts Used:** Visualization, Outlier Detection.

**Code:**

```
plt.figure(figsize=(12,6))
```

```
sns.boxplot(x='drg_definition', y='average_total_payments', data=df)
```

```
plt.xticks(rotation=90)
```

```
plt.title("Hospital Charges by Department (DRG)")
```

```
plt.xlabel("Department (DRG Definition)")
```

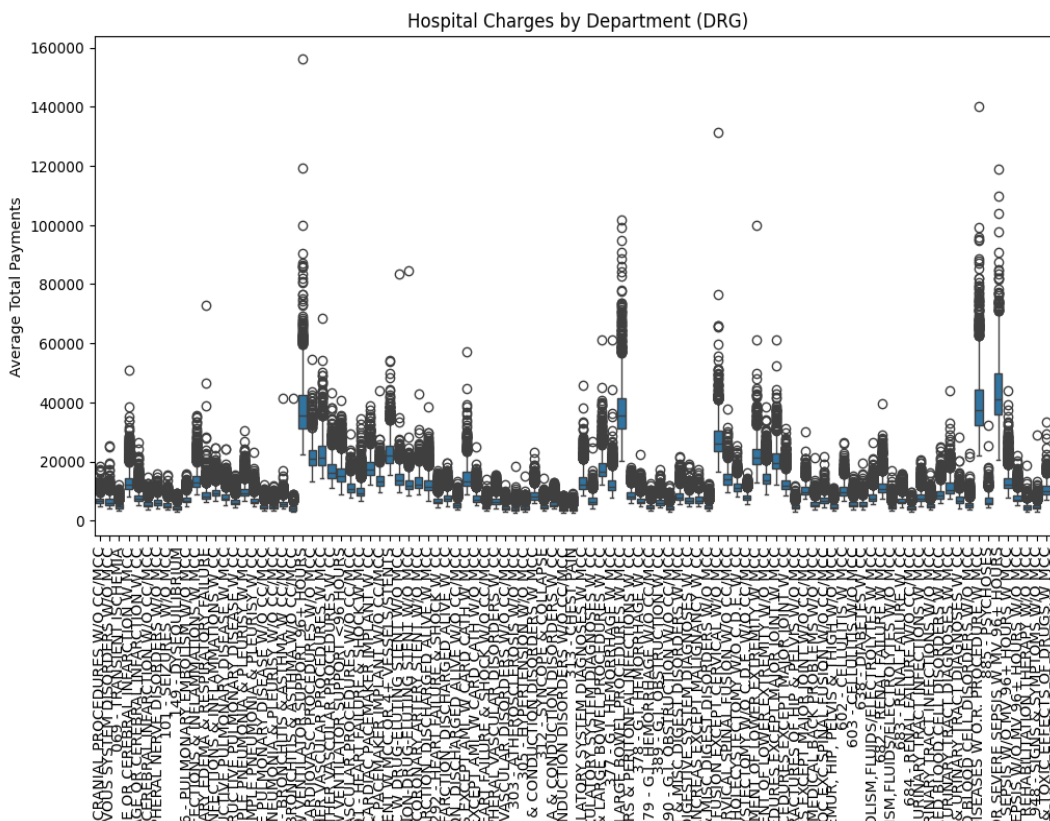
```
plt.ylabel("Average Total Payments")
```

```
plt.show()
```

### Expected Output:

- Boxplot showing distribution of charges across different DRGs.

## Actual Output:



Activate Windows  
Go to Settings to activate Windows.

### Data Engineering Insight:

Boxplots quickly highlight outliers (extremely high charges), useful for anomaly detection.

## Exercise 18: Pie chart – patient gender breakdown

**Purpose:** Show categorical distribution of gender.

**Concepts Used:** Visualization, Profiling.

**Code:**

```
gender_counts = df['gender'].value_counts()
```

```
plt.figure(figsize=(6,6))
```

```
plt.pie(gender_counts, labels=gender_counts.index, autopct='%1.1f%%',  
startangle=90)
```


```
plt.title("Patient Gender Breakdown")
```

```
plt.show()
```

**Expected Output:**

- Pie chart with male/female proportions.

**Actual Output:**



---

**Data Engineering Insight:**

Pie charts provide quick insights into categorical balance, useful for demographic profiling.

## Exercise 19: Line plot – monthly patient admissions

**Purpose:** Show time-series trends in patient admissions.

**Concepts Used:** Visualization, Time-Series Analysis.

**Code:**

```
monthly_admissions = df.groupby(df['admission_date'].dt.to_period("M")).size()
```

```
plt.figure(figsize=(10,5))
```

```
monthly_admissions.plot(kind='line', marker='o')
```

```
plt.title("Monthly Patient Admissions")
```

```
plt.xlabel("Month")
```


```
plt.ylabel("Number of Admissions")
```

```
plt.show()
```

**Expected Output:**

- Line plot showing trends in admissions per month.

**Actual Output:**



**Data Engineering Insight:**

Line plots help validate seasonality or irregular spikes in patient admissions.

## Exercise 20: Heatmap – correlation validation

**Purpose:** Visually validate correlation between numerical variables.

**Concepts Used:** Visualization, Correlation, Validation.


**Code:**

```
plt.figure(figsize=(8,6))  
  
sns.heatmap(corr_matrix, annot=True, cmap="coolwarm", fmt=".2f")  
  
plt.title("Correlation Heatmap")  
  
plt.show()
```

**Expected Output:**

- Heatmap with correlations (e.g., charges vs age vs stay).

**Actual Output:**



**Data Engineering Insight:**

Heatmaps give a visual summary of correlations, making it easier to spot strong/weak relationships.